

## Purpose

This collection consists of:

Component	Purpose
tJobInstanceStart	Register a job run and provide information about the previous job run. Setup the logging facility.
tJobInstanceEnd	Updates the job run data with metrics, job results and error information.
tJobDataRangeScanner	Collects min/max values of timestamps or numeric values within a data flow. These min/max values for timestamps and any other values are set for the job-time-range and job-value-range.
tJobInstanceLiveCheck	Checks the entries of the job run registry for dead or broken job instances and cleans up the job registry.

These components help to track the execution of jobs in a database table.

Advantages of these components:

- Provides a unique numeric id for the job to mark all data sets processed by the job
- Start-/Stop timestamps
- Return code and error messages (collects all messages)
- Host and PID of the process running this job
- Supports incremental loading
- Supports restart capabilities by providing information about the previous job run
- Key figures about moved data sets
- Snapshot of the context at the start and at the end of the job
- Detects the minimum and maximum of value for a flow
- Tracks the memory usage and detects peaks
- Detects parallel job runs (also based on the work item) and provide steering possibilities

## Talend-Integration

This component can be found in the palette under Management

This component provides several return values.

## Component tJobInstanceStart

### Basic Settings

Property	Content
Use data source (connection pool)	If true, the component takes the database connection from a database connection pool.
Data source alias	The name of the database connection pool providing the database connection.
Database Connection	Any database connection pointing to the schema with the control tables. The main table is JOB_INSTANCES holding all key information.
Job Name	Name of the job. The default is using the build-in variable jobName
Task name (Job Display Name)	To distinguish job executions of the same job for different purposes, you can provide an additional declaration of the current task (mostly relevant for more generic jobs, that are called for different tasks). Best practice is to extend the job name with "-" or "#" and add a name extension to declare this way the task name. This information will be stored in the

	field job_display_name.
Process Instance Name	Since multiple jobs / job instances might be executed within the same process (via tRunJob) this property can be set to a more general job name, representing the whole process.
Job Work Item	Text describing the work item (e.g. a file name or a source object name) processed by this job instance. It can be used by later executions to determine already processed work items or to track individual increment boundaries for multiple source instances.
Take empty as null	If the value for work item is an empty string, it will be stored as null
Time range start	If the job has to process data selected by a time range, declare the start of the timerange here.
Time range end	If the job has to process data selected by a time range, declare the end of the timerange here.
Value range start	If the job has to process a partition of data selected by an id range or any other value ranges, you can store the start of the requested range here.
Value range end	If the job has to process a partition of data selected by an id range or any other value ranges, you can store the end of the requested range here.
Write Job instance ID to	To inject the job instance id into the processd data, typically a context variable will be used. Provide the context variable to be loaded with the job instance id here.
Singleton Behavior	See the section Singleton Behavior later in this document.
Return last instance result	Fetches the information about the last run of this job. All information available as return values of the tJobInstanceStart component.
Last successful	The last run must have been a successful run of this job (all others will be ignored)
Last run must have inserted or deleted data	The last run must have data inserted or deleted. This will be detected via the couter columns. See the properties of tJobInstanceEnd.
Last run must have read data	The last run must have read data to be relevant as choice for the last run. This will be detected via the couter columns. See the properties of tJobInstanceEnd.
For the current work item	If true: only jobs with the same work item will be used to get the last instance results. This allows you to have one job for multiple purposes with the full support of the “Return last instance result” feature.
Collecting job instances ids running after previous run	Returns as comma separated list all instance ids of all job, which was running after the last run of this job. This helps to implement incremental jobs. It is necessary to write the job instance id into every data set proceed by the job.
Source job names	Filter the jobs which should part of the list above. This helps to keep the list small in case of having a lot of unrelated jobs in the system.
OK Result Codes	This is a String containing a comma-separated list of all return codes, which are related to a successful run. If you want using different return codes for OK please take care the tRunJob components does not die.
Replacement for prev...	These values can be used to simplify the job design because you can avoid the check if the job has a previous run. Please refer to the return values.
...-job-instance-id	Set here a default long value for the return value PREV_JOB_INSTANCE_ID
...-job-start-date	Set here the Date typed default value for a previous job start date
...-time-range-end	Set here the Date typed default value for a previous time range end.
...-value-range-end	Set here the String typed default value for a previous value range end.
...-result-item	Set here the String typed default value for a previous job result item.
Set UTC as default time zone	This changes the default setting of the virtual machine for time zone from the local time zone to UTC. It affects the current JVM instance (=also jobs called by tRunJob that are not an independent child jobs).
Memory Usage Monitoring	This starts a thread, which collects every second, the used memory and detect the maximum and when it happened. Values are provided in global map after tJobInstanceEnd has finished.
Print job instance id to the console	If true the “Expresion to pring” will be printed to System.out

Expression to print	Define here how to print the job instance id to system.out. You can add any available data as you like. It is recommended to print at least the jobname and the job instance id
---------------------	---

## Advanced Settings

Max message length	
Generate Job Instance Id	(experimental) This generates the Job Instance Id from systemtime and a large random number. The generated number should be unique even, when multiple instances are started in the same second.
Job instance ID is auto increment	This have to be switched on if the table use an auto increment for the job_instance_id column e.g. this is supposed for MySQL.
Sequence expression	The valid SQL Expression to get an unused job_instance_id from the database. In general this is a "nextval" of the DB sequence. Syntax depends on the database. Examples: Oracle:        job_instance_id_seq.nextval PostgreSQL:   nextval('dwh_manage.job_instance_id_seq') DB2:    NEXTVAL FOR dwh_manage.job_instance_id_seq
Reduce resume logs	
Enable JMX	(Experimental JMX Connector to retrieve current values via JMX)
tStatCatcher Statistik	Enable trace output to talend statcatcher

## Return values of tJobInstanceStart

Return value	Content
ERROR_MESSAGE	Last error message. Unfortunately this is not the error message from the actually running job. This message is build from the tRunTask component. The current TAC web service does not provide this message.
JOB_INSTANCE_ID	The job instance id used for this job run.
SOURCE_JOB_INSTANCE_ID_LIST	List of all job instance ids which are executed after the last run if this job. This way it is possible to implement incremental steering. The list can easily be used in SQL e.g. : ...where job_instance_id in (" + ((String)globalMap.get("tJobInstanceStart_1_SOURCE_JOB_INSTANCE_ID_LIST") + ") .....")
JOB_START_DATE	The start date of the current job run.
PREV_JOB_EXISTS	If true means the job was running in the past at least one time.
PREV_JOB_START_DATE	If a previous job run exists (otherwise null): Contains the start date of the previous job
PREV_JOB_STOP_DATE	If a previous job run exists (otherwise null): Contains the stop date of the previous job
PREV_JOB_INSTANCE_ID	If a previous job run exists (otherwise null): Contains the ID of the previous job
PREV_JOB_TALEND_PID	If a previous job run exists (otherwise null): Contains the Talend-PID of the previous job
PREV_JOB_HOST_PID	If a previous job run exists (otherwise null): Contains the Host-PID (means the process ID of the operating system for this JVM) of the previous job
PREV_JOB_HOST_NAME	If a previous job run exists (otherwise null):

	Contains the name of the host where the previous job was running
PREV_TIME_RANGE_START	If a previous job run exists (otherwise null): Contains the time range start of the previous job
PREV_TIME_RANGE_END	If a previous job run exists (otherwise null): Contains the time range end of the previous job
PREV_VALUE_RANGE_START	If a previous job run exists (otherwise null): Contains the value range start of the previous job
PREV_VALUE_RANGE_END	If a previous job run exists (otherwise null): Contains the value range end of the previous job
PREV_JOB_RETURN_CODE	If a previous job run exists (otherwise null): Contains the return code of the previous job
PREV_WORK_ITEM	If a previous job run exists (otherwise null): Contains the previous work item of the previous job
PREV_RESULT_ITEM	If a previous job run exists (otherwise null): Contains the result item of the previous job
PREV_COUNT_INPUT	If a previous job run exists (otherwise null): Contains the count inserts of the previous job
PREV_COUNT_OUTPUT	If a previous job run exists (otherwise null): Contains the count outputs of the previous job
PREV_COUNT_UPDATED	If a previous job run exists (otherwise null): Contains the count updates of the previous job
PREV_COUNT_DELETED	If a previous job run exists (otherwise null): Contains the count deletes of the previous job
PREV_COUNT_REJECTS	If a previous job run exists (otherwise null): Contains the count rejects of the previous job

## Singleton Behavior

This component has the ability based of the information of the JOB\_INSTANCE\_STATUS table to detect already running instances of the same job (optional with the same work item).

The necessary basic settings to use this feature are:

Property	Content
Check if another job instance of the same job is already running	If true, the component checks the JOB_INSTANCE_STATUS table for a already running job instance
Singleton for the work item	Use the work item to identify a job instance.
Prevent creating a job instance status entry if another job instance is already running	The option helps to avoid useless entries in the JOB_INSTANCE_STATUS table.

Dedicated return values

Return value	Content
JOB_RUNS_ALONE	The result of the singleton-check. This variable is not set if the check has not happened, this should help to prevent using this value if the actual check was not performed.
ALREADY_RUNNING_JOB_START_DATE	Start date (+time) of the already running job
ALREADY_RUNNING_JOB_INSTANCE_ID	The job instance id of the other job still running
ALREADY_RUNNING_JOB_HOST_NAME	On which server the other job instance is still running
ALREADY_RUNNING_JOB_HOST_PID	The native process id of the other job instance
ALREADY_RUNNING_JOB_WORK_ITEM	What work item the other job is still processing

## Component tJobInstanceEnd

### Basic Settings

Property	Content
Use separate connection	If true: the end component uses a separate database connection. This could help for jobs running very long to avoid problems with long standing connections.
Connection	Choose here the separate connection if you have chosen the option above. It must be a different connection component as used for the start component.
Job Instance Start Component	Choose here the tJobInstanceStart component. Both components depend on each other.
Job Result	A string representation of the result of the current job. In case the job <u>creates</u> a file, it is a good idea to put here the file path.
Time range start	If the job has processed data of a defined time interval, you can store the start of the interval here.
Time range end	If the job has processed data of a defined time interval, you can store the end of the interval here. This can be used as a starting point for the next iteration.
Value range start	If the job has process a partition of data selected by an id range or any other value ranges, you can store the start of the retrieved/requested range here
Value range end	See Value range start. This is the end of the retrieved/requested range.
Save named counters	Counters can be named; in this case the counter value will additionally be inserted in the table JOB_INSTANCE_COUNTERS
Delete previous successful job instances by work item	If checked, the component deletes all successful previous job instances with the same work item. This helps in case of the table JOB_INSTANCE_STATUS will be used to keep track of the current data in the DWH and repeated job runs with the same work item replaces previous data.
Close Connection	Closes the connection used for managing the job registration
Input Counters	Counters describing the result of the job can be added here. The sum of all counters will be written in the JOB_INSTANCE_STATUS table in COUNT_INPUTS. By switching off the flag “Add” you can subtract a value instead of adding it. The name column provides the name to be used in the JOB_INSTANCE_COUNTERS table (see Save named counters option)
Output Counters	See Input Counters. Will be used for column COUNT_OUTPUTS
Update Counters	See Input Counters. Will be used for column COUNT_UPDATED
Reject Counters	See Input Counters. Will be used for column COUNT_REJECTED
Delete Counters	See Input Counters. Will be used for column COUNT_DELETED

As Counter typically the NB\_LINE return values of the input or output components can be used. In case of the job has more the one output it is recommended to set names for particular counters to keep the distinct counter values.

## Return values of tJobInstanceEnd

Return value	Content
ERROR_MESSAGE	Last error message.
RETURN_CODE	The retrieved return code of the current job
RETURN_MESSAGE	The created return messages. This message contains all error messages from all components throwing an error.
MEMORY_AVAILABLE	This is the memory (in byte) what is maximum available for the job. Typically it is set with the JVM parameter <code>-Xmx1024m</code> (e.g. for 1GB RAM)
MEMORY_MAX_USED	The maximum of the used memory (in byte) what was allocated in the JVM. Please keep in mind, if you call other jobs with tRunJob (not independently) they must be taken into account because they use the same JVM instance.
MEMORY_MAX_USED_PERCENTAGE	The percentage between the available memory and the maximum used memory as value between 0 and 100.

## Component tJobDataRangeScanner

This component measures the min/max values from any values (like IDs) or timestamps (like last\_modified). These min/max values will be used to set the column values in the table job\_instance\_status (value\_range\_start, value\_range\_end, time\_range\_start, time\_range\_end).

All these values can be retrieved for the last run as return values of the tJobInstanceStart component.

Please refer to the return values of the tJobInstanceStart component.

### Basic settings

Property	Content
Job Instance Start Component	Choose here the tJobInstanceStart component. Both components depend on each other.
Schema	This is necessary to have the schema column available. It is not supposed to change anything here
Configure Extraction	For every schema column you can define for which range it will be checked: Time range or Value range. The min and max values will be found even the component runs in iteration.

### Return values

Return value	Content
ERROR_MESSAGE	Last error message in case of the range detection fails for a column.
TIME_RANGE_START	The min value for the measured time range.
TIME_RANGE_END	The max value for the measured time range.
VALUE_RANGE_START	The min value for the measured value range as Long or String
VALUE_RANGE_END	The max value for the measured value range as Long or String
NB_LINE_AGGREGATED	The number of rows for this component measured over all iterations



## Component tJobInstanceLiveCheck

This component checks if jobs still alive. To do this, you have to build a very simple job (checkout the scenarios) and let them run on every job server you have.

### Basic settings

Property	Content
Database Connection	Any database connection pointing to the schema with the control tables. The main table is JOB_INSTANCE_STATUS holding all key information.
Close Connection	If true the connection will be closed at the end of the component processing
Schema	This component provides an input flow providing information about cleaned job instances
Last system start	If the last system start could be determined (currently there is not platform independent implementation to get this information automatically) all older job instance starts will be cleaned.

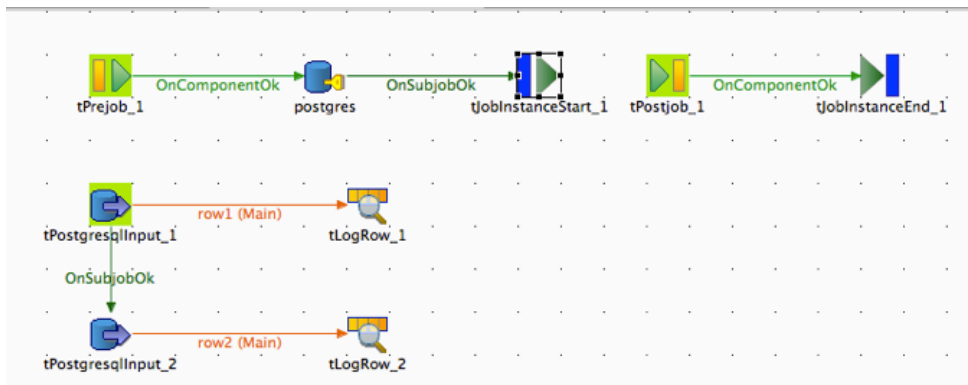
### Return values

Return value	Content
ERROR_MESSAGE	Error message if something in the processing of the component it self went wrong
COUNT_RUNNING_PROCESSES	The number of all running processes on the current server (regardless if this is a Talend job or not)
COUNT_RUNNING_JOB_INSTANCES	The number of as running declared job instances
COUNT_BROKEN_JOB_INSTANCES	The number of recognized broken job instances
NB_LINE	Number of rows in the data input flow

The schema is fully commented and provides the values of the JOB\_INSTANCE\_STATUS table for the broken instances.

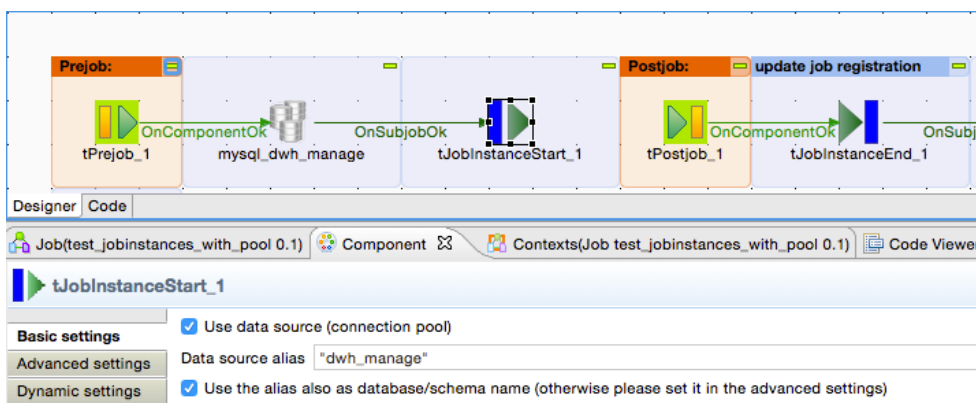
## Scenarios

### Scenario 1: Simple Job monitoring



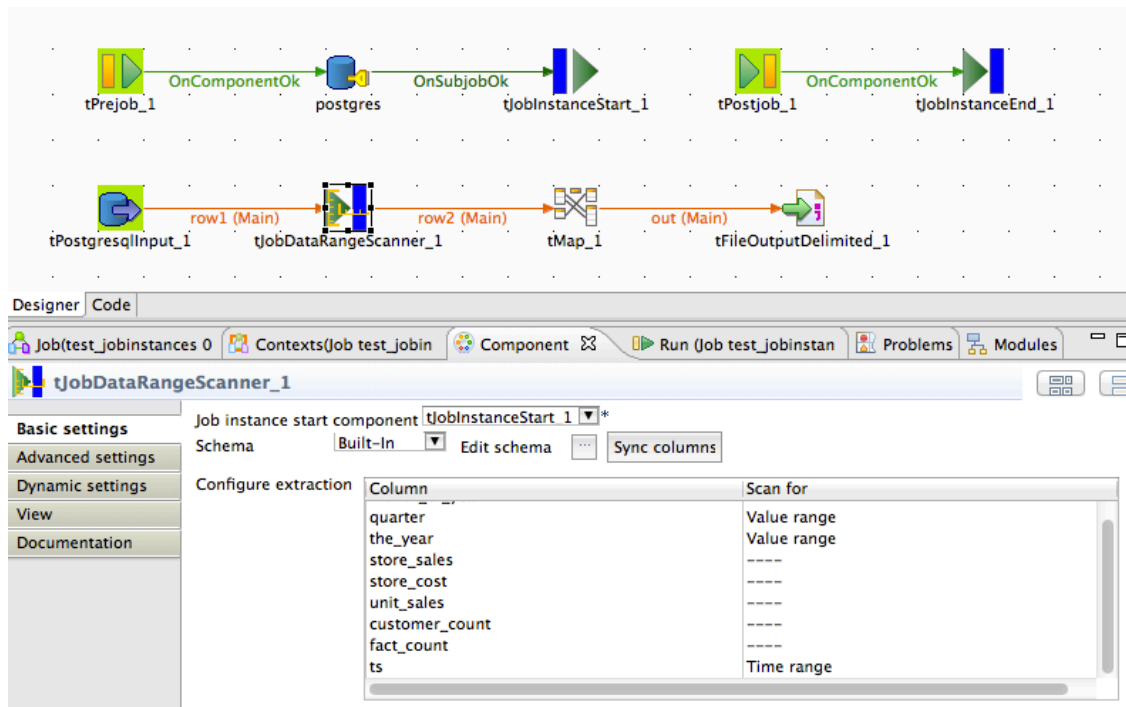
The typical usage is to use **tPrejob** component to trigger the **tJobInstanceStart** component and the **tPostjob** component to trigger **tJobInstanceEnd** component.

### Scenario 2: Using a connection pool



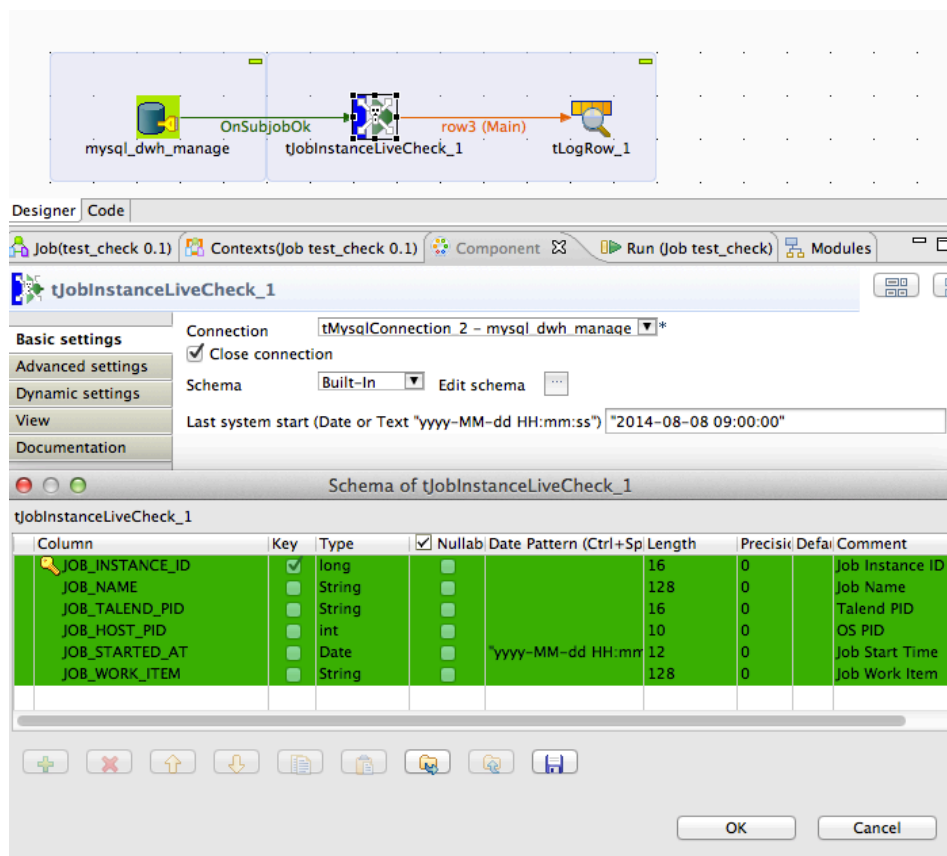
The connection pool can be established anywhere unless it is before the initialization of the **tJobInstanceStart** component. Also, child jobs can use the same connection pool. It is a good practice to name the pool like the addressed database. This way the configuration more convenient.

### Scenario 3: Measure the time ranges and/or value ranges



In this Scenario the flow will be scanned for the start and end values for a time range and the value range. These values could be used to ensure the job quality or to start the next run from the previous end.

Scenario for tJobInstanceLiveCheck



This example job shows the main purpose of the component. Such kind of job has to run frequently on every job server (servers on which the jobs run).

The component set for broken job instances the return code=999 and as return message “Process died”. This information can be used to clean up all depending data structures.

## Scenario 4: Using the singleton check

The check depends in the information of the job\_instance\_status table. If you decide to use this check it is highly recommended to establish also a live check for job instances (use tJobInstanceLiveCheck component).

The screenshot displays the SAP Job Designer interface. At the top, a job flow diagram shows the sequence of components: Prejob, tPrejob\_1, postgres\_dwh\_manage, tJobInstanceStart\_1, tPostjob\_1, and tJobInstanceEnd\_1. Below the diagram, the configuration panel for the selected component, tJobInstanceStart\_1, is visible. The panel includes sections for Basic settings, Singleton Behavior, and Information about previous job runs.

**Basic settings**

- ☐ Use data source (connection pool)
- Connection: tPostgresolConnection\_1 - postgres\_dwh\_manage
- ☐ Close connection (Depends on using a different one for tJobInstanceEnd!)
- Job name: jobName
- Task name (if the job runs for different tasks): jobName
- Process instance name:
- Job work item: "workitem3"
- ☐ Take empty value as null
- Time range start:
- Time range end:
- Value range start:
- Value range end:
- Context variable to be filled with Job-Instance-ID (Long-typed):

**Singleton Behavior**

- ☒ Check if another job instance of the same job is already running
- ☒ Singleton for the work item (expects not null work item!)

Note: You can use the an if trigger with the return value JOB\_RUNS\_ALONE (boolean) to trigger the process in case no other instance is running or let the process die by

- ☒ Prevent creating a job-instance-status entry (and log entries) if another job instance is already running.

Note: You can avoid a register entry for a job which should actually not run.

**Information about previous job runs**

- ☒ Return last instance results
- ☐ Last successful
- ☒ Last run must have inserted, updated or deleted data
- ☐ Last run must have read data
- ☒ For the current work item
- ☐ Only within the same root job (Depends on you have a root job equipped with theses components)
- ☐ Collect Job Instance IDs running after previous run

The if-trigger contains the return variable JOB\_RUNS\_ALONE to trigger the actual job or in this case to let the job die. The tWarn component is an very simple example of a job. Start here with your actual work.

## Checklist to use this component suite

1. Use an OLTP database for the tables used by this component. Column oriented database are mostly too slow for the possible high frequency insert/updates from a large number of simultaneous running jobs! Typical bad choices: Teradata, Infobright engine of MySQL... Typical capable databases: Oracle, MySQL, DB2, PostgreSQL, H2
2. Build a separate database schema (e.g. call it dwh\_manage or dwh\_meta or what you want ;-)
3. Think about the primary key in the JOB\_INSTANCE\_STATUS table. You can use a database sequence (this is the preferred way) or you can use a self-incrementing data type (e.g. serial or identity column types). If you use a sequence, please set the SQL code to get the next value in the advanced settings of the tJobInstanceStart component. By the way, in large projects it could be helpful to set these settings per default directly in the component in the tJobInstanceStart\_java.xml file of the component.
4. If you use the component, please check if the job is long running job and, in this case, use a separate connection for the tJobInstanceEnd component or consider the usage of a connection pool to avoid problems with server-side disconnected database connections.

Table: job\_instance\_counters

column name	column type	description
job_instance_id	bigint	Related job_instance_id
counter_name	varchar(128)	Name of the counter. Set this name in the tJobInstanceEnd as name for an entry in the counter table.
counter_type	Varchar(20)	Type of the counter (input, output,...)
counter_value	integer	The value of the counter

## Create table scripts for the tables

Not all databases are capable to work for this use case. Generally, all OLTP databases work fine. It could be problematic to host these tables on a column-based database or a database with distributed storage like Teradata. Such database tends to be very slow for single inserts and updates or they lock the whole table while such operations and this could lead to significant performance problems!

The well-tested databases are MySQL (MyISAM or INNODB), Oracle, PostgreSQL, H2 and IBM DB2

In case of MySQL it is recommended using a serial data type for the column

JOB\_INSTANCE\_STATUS(JOB\_INSTANCE\_ID).

In the advanced settings of the tJobInstanceStart component it is possible to declare the schema and the table names.

The option Job Instance ID is auto increment allows the usage of auto increment column for JOB\_INSTANCE\_ID in the table JOB\_INSTANCE\_STATUS. In former releases some tables had slightly different names, but the meaning and structure is the same. You can adapt old names in the configuration if the tJobInstanceStart advanced settings.

## MySQL (use auto increment)

```
CREATE TABLE JOB_INSTANCE_STATUS (
  JOB_INSTANCE_ID BIGINT UNSIGNED NOT NULL AUTO_INCREMENT,
  PROCESS_INSTANCE_ID BIGINT DEFAULT NULL,
  PROCESS_INSTANCE_NAME VARCHAR(255) DEFAULT NULL,
  JOB_NAME VARCHAR(255) NOT NULL,
  JOB_PROJECT varchar(128),
  JOB_DISPLAY_NAME VARCHAR(255) DEFAULT NULL,
  JOB_GUID VARCHAR(100) NOT NULL,
  JOB_EXT_ID VARCHAR(255) DEFAULT NULL,
  JOB_INFO VARCHAR(255) DEFAULT NULL,
  ROOT_JOB_GUID VARCHAR(100) DEFAULT NULL,
  WORK_ITEM VARCHAR(1024) DEFAULT NULL,
  TIME_RANGE_START TIMESTAMP(3) NULL DEFAULT NULL,
  TIME_RANGE_END TIMESTAMP(3) NULL DEFAULT NULL,
  VALUE_RANGE_START VARCHAR(512) DEFAULT NULL,
  VALUE_RANGE_END VARCHAR(512) DEFAULT NULL,
  JOB_STARTED_AT TIMESTAMP(3) NULL DEFAULT NULL,
  JOB_ENDED_AT TIMESTAMP(3) NULL DEFAULT NULL,
  JOB_RESULT VARCHAR(1024) DEFAULT NULL,
  COUNT_INPUT INT(11) DEFAULT NULL,
  COUNT_OUTPUT INT(11) DEFAULT NULL,
  COUNT_UPDATED INT(11) DEFAULT NULL,
  COUNT_REJECTED INT(11) DEFAULT NULL,
  COUNT_DELETED INT(11) DEFAULT NULL,
  RETURN_CODE INT(11) DEFAULT NULL,
  RETURN_MESSAGE TEXT,
  HOST_NAME VARCHAR(255) DEFAULT NULL,
  HOST_PID INT(11) DEFAULT NULL,
  HOST_USER VARCHAR(128) DEFAULT NULL,
  PRIMARY KEY (JOB_INSTANCE_ID)
) DEFAULT CHARSET=UTF8;

CREATE INDEX JOB_INSTANCE_STATUS_JOB_GUID ON JOB_INSTANCE_STATUS(JOB_GUID);
CREATE INDEX JOB_INSTANCE_STATUS_JOB_NAME ON JOB_INSTANCE_STATUS(JOB_NAME);

CREATE TABLE JOB_INSTANCE_COUNTERS (
  JOB_INSTANCE_ID BIGINT NOT NULL,          -- reference to the job instance
  COUNTER_NAME VARCHAR(128) NOT NULL,       -- name of the counter set in tJobInstanceEnd for a counter
  COUNTER_TYPE VARCHAR(20),                -- type of the counter
  COUNTER_VALUE INTEGER,                   -- value of the counter
  CONSTRAINT PK_JOB_INSTANCE_COUNTERS PRIMARY KEY (JOB_INSTANCE_ID, COUNTER_NAME));
```

## PostgreSQL (uses a sequence)

```
--drop table dwh_manage.job_instances;
create table dwh_manage.job_instance_status (
    job_instance_id bigint not null,
    process_instance_id bigint,
    process_instance_name varchar(255),
    job_name varchar(255) not null,
    job_project varchar(128),
    job_info varchar(512),
    job_display_name varchar(255),
    job_guid varchar(100) not null,
    job_ext_id varchar(255),
    root_job_guid varchar(100),
    work_item varchar(1024),
    time_range_start timestamp,
    time_range_end timestamp,
    value_range_start varchar(512),
    value_range_end varchar(512),
    job_started_at timestamp not null,
    job_ended_at timestamp,
    job_result varchar(1024),
    count_input integer,
    count_output integer,
    count_updated integer,
    count_rejected integer,
    count_deleted integer,
    return_code integer,
    return_message varchar(1024),
    host_name varchar(255),
    host_pid integer,
    host_user varchar(128),
    constraint job_instances_pkey primary key (job_instance_id));

create index job_instances_job_guid on dwh_manage.job_instance_status(job_guid);
create index job_instances_job_name on dwh_manage.job_instance_status(job_name);

--drop sequence dwh_manage.seq_job_instance_id;
create sequence dwh_manage.seq_job_instance_id start with 1;

--drop table dwh_manage.job_instance_counters;
create table dwh_manage.job_instance_counters (
    job_instance_id bigint not null,
    counter_name varchar(128) not null,
    counter_type varchar(20),
    counter_value integer not null);

create index job_instance_counters_idx on dwh_manage.job_instance_counters(job_instance_id,
counter_name);
```



## Oracle (uses a sequence)

```
-- drop table JOB_INSTANCE_STATUS;
CREATE TABLE JOB_INSTANCE_STATUS (
  JOB_INSTANCE_ID NUMBER(19) NOT NULL,
  PROCESS_INSTANCE_ID NUMBER(19),
  PROCESS_INSTANCE_NAME VARCHAR2(255),
  JOB_NAME VARCHAR2(255) NOT NULL,
  JOB_PROJECT VARCHAR2(128),
  JOB_INFO VARCHAR2(512),
  JOB_DISPLAY_NAME VARCHAR2(255),
  JOB_GUID VARCHAR2(100) NOT NULL,
  JOB_EXT_ID VARCHAR2(255),
  ROOT_JOB_GUID VARCHAR2(100),
  WORK_ITEM VARCHAR2(1024),
  TIME_RANGE_START TIMESTAMP,
  TIME_RANGE_END TIMESTAMP,
  VALUE_RANGE_START VARCHAR2(512),
  VALUE_RANGE_END VARCHAR2(512),
  JOB_STARTED_AT TIMESTAMP NOT NULL,
  JOB_ENDED_AT TIMESTAMP,
  JOB_RESULT VARCHAR2(1024),
  COUNT_INPUT INTEGER,
  COUNT_OUTPUT INTEGER,
  COUNT_UPDATED INTEGER,
  COUNT_REJECTED INTEGER,
  COUNT_DELETED INTEGER,
  RETURN_CODE INTEGER,
  RETURN_MESSAGE VARCHAR2(1024),
  HOST_NAME VARCHAR2(255),
  HOST_PID INTEGER,
  HOST_USER VARCHAR(128),
  CONSTRAINT JOB_INSTANCES_PKEY PRIMARY KEY (JOB_INSTANCE_ID));

CREATE INDEX JOB_INSTANCES_JOB_GUID ON JOB_INSTANCE_STATUS(JOB_GUID);
CREATE INDEX JOB_INSTANCES_JOB_NAME ON JOB_INSTANCE_STATUS(JOB_NAME);

CREATE SEQUENCE SEQ_JOB_INSTANCE_ID START WITH 1;

CREATE TABLE JOB_INSTANCE_COUNTERS (
  JOB_INSTANCE_ID NUMBER(19) NOT NULL,
  COUNTER_NAME VARCHAR2(128) NOT NULL,
  COUNTER_TYPE VARCHAR2(20),
  COUNTER_VALUE INTEGER NOT NULL);

CREATE INDEX JOB_INSTANCE_COUNTERS_IDX ON JOB_INSTANCE_COUNTERS(JOB_INSTANCE_ID, COUNTER_NAME);
```

## IBM DB2

```
--drop table dwh_manage.job_instance_status;
create table dwh_manage.job_instance_status (
    job_instance_id bigint not null,
    process_instance_id bigint,
    process_instance_name varchar(255),
    job_name varchar(255) not null,
    job_project varchar(128),
    job_info varchar(512),
    job_display_name varchar(255),
    job_guid varchar(100) not null,
    job_ext_id varchar(255),
    root_job_guid varchar(100),
    work_item varchar(1024),
    time_range_start timestamp,
    time_range_end timestamp,
    value_range_start varchar(512),
    value_range_end varchar(512),
    job_started_at timestamp not null,
    job_ended_at timestamp,
    job_result varchar(1024),
    count_input integer,
    count_output integer,
    count_updated integer,
    count_rejected integer,
    count_deleted integer,
    return_code integer,
    return_message varchar(1024),
    host_name varchar(255),
    host_pid integer,
    host_user varchar(128),
    constraint job_instances_pkey primary key (job_instance_id));

create index dwh_manage.job_instances_job_guid on dwh_manage.job_instance_status(job_guid);
create index dwh_manage.job_instances_job_name on dwh_manage.job_instance_status(job_name);

--drop sequence dwh_manage.job_instance_id;
create sequence dwh_manage.seq_job_instance_id start with 1;

--drop table dwh_manage.job_instance_counters;
create table dwh_manage.job_instance_counters (
    job_instance_id bigint not null,
    counter_name varchar(128) not null,
    counter_type varchar(20),
    counter_value integer not null);

create index job_instance_counters_idx on dwh_manage.job_instance_counters(job_instance_id,
counter_name);
```

## Exasol

```
-- drop table dwh_manage.job_instance_status;
create table dwh_manage.job_instance_status (
  job_instance_id bigint identity primary key,
  process_instance_id bigint,
  process_instance_name varchar(255),
  job_name varchar(255) not null,
  job_project varchar(512) UTF8,
  job_info varchar(512) UTF8,
  job_display_name varchar(255) UTF8,
  job_guid varchar(100) UTF8 not null,
  job_ext_id varchar(255) UTF8,
  root_job_guid varchar(100) UTF8,
  work_item varchar(1024) UTF8,
  time_range_start timestamp,
  time_range_end timestamp,
  value_range_start varchar(512) UTF8,
  value_range_end varchar(512) UTF8,
  job_started_at timestamp not null,
  job_ended_at timestamp,
  job_result varchar(1024) UTF8,
  count_input integer,
  count_output integer,
  count_updated integer,
  count_rejected integer,
  count_deleted integer,
  return_code integer,
  return_message varchar(4000) UTF8,
  host_name varchar(255) UTF8,
  host_pid integer,
  host_user varchar(128) UTF8);

create table dwh_manage.job_instance_counters (
  job_instance_id bigint not null,
  counter_name varchar(128) not null,
  counter_type varchar(20),
  counter_value integer not null);
```